

ORACLE®

Certified Professional




Oracle Certified  
Professional Java  
Programmer

Microsoft Certified  
Professional

sohailimran@yahoo.com



سہیل عمران Sohail IMRAN



*intro*

# Common systems we run today



Cache



OLTP



Real-time  
data processing



NoSQL/  
NewSQL



Hadoop



OLAP/  
EDW



Messaging



In-memory  
processing

The challenges that supporting all these systems brings:

- They have to build the in-house knowledge to administer and keep the systems up and healthy. Systems fail, especially complicated distributed systems, and being open source, many of these systems don't have the tooling to facilitate easy management.
- Data exchange between systems is painful, primarily due to the volume of data and the lack of tooling for the data movement. Large, expensive projects ensue.
- Each system has to solve the same distributed problems, such as fault tolerance, distributed storage, log handling, and resource scheduling.
- <https://www.edureka.co/blog/mapreduce-tutorial/>

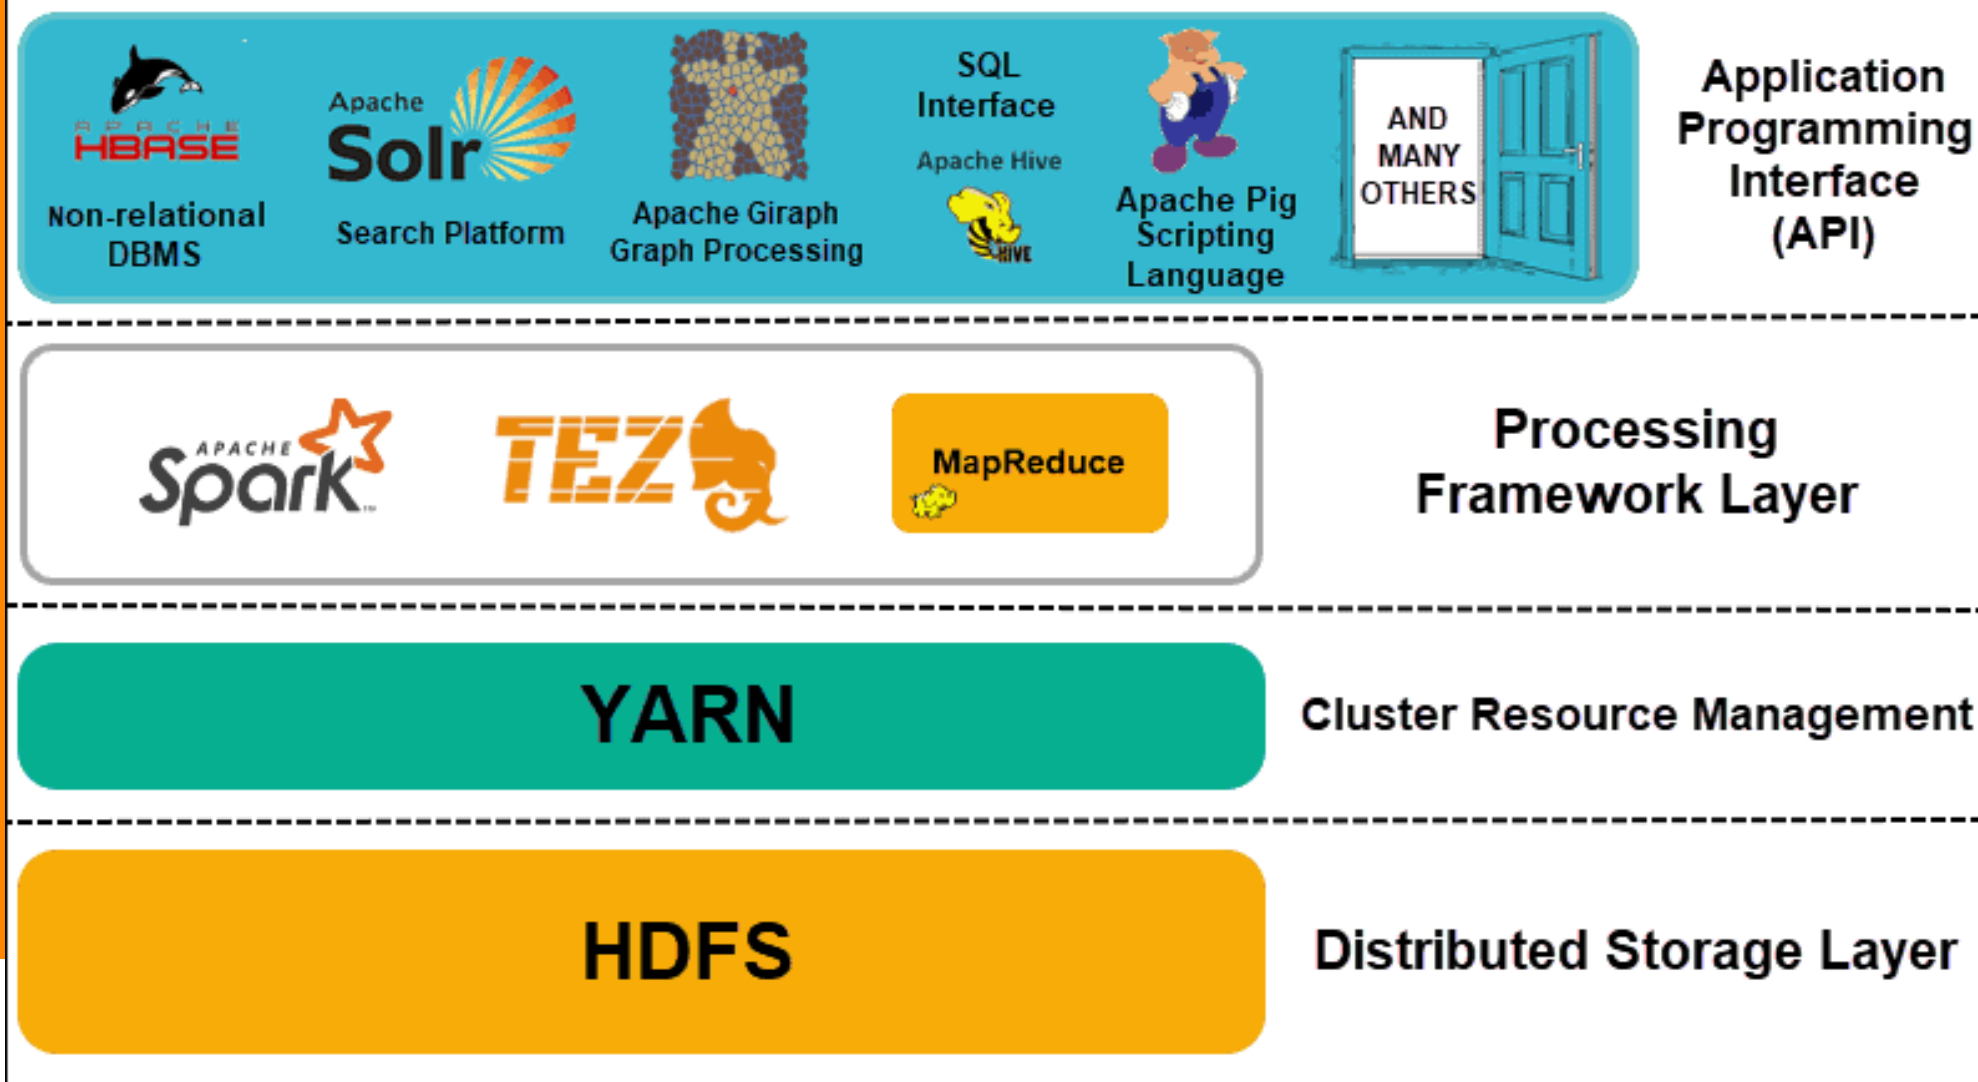
# a distributed framework for big data

- ▶ Apache top level project, open-source implementation of frameworks for reliable, scalable, distributed computing and data storage.
- ▶ It is a flexible and highly-available architecture for large scale computation and data processing on a network of commodity hardware.

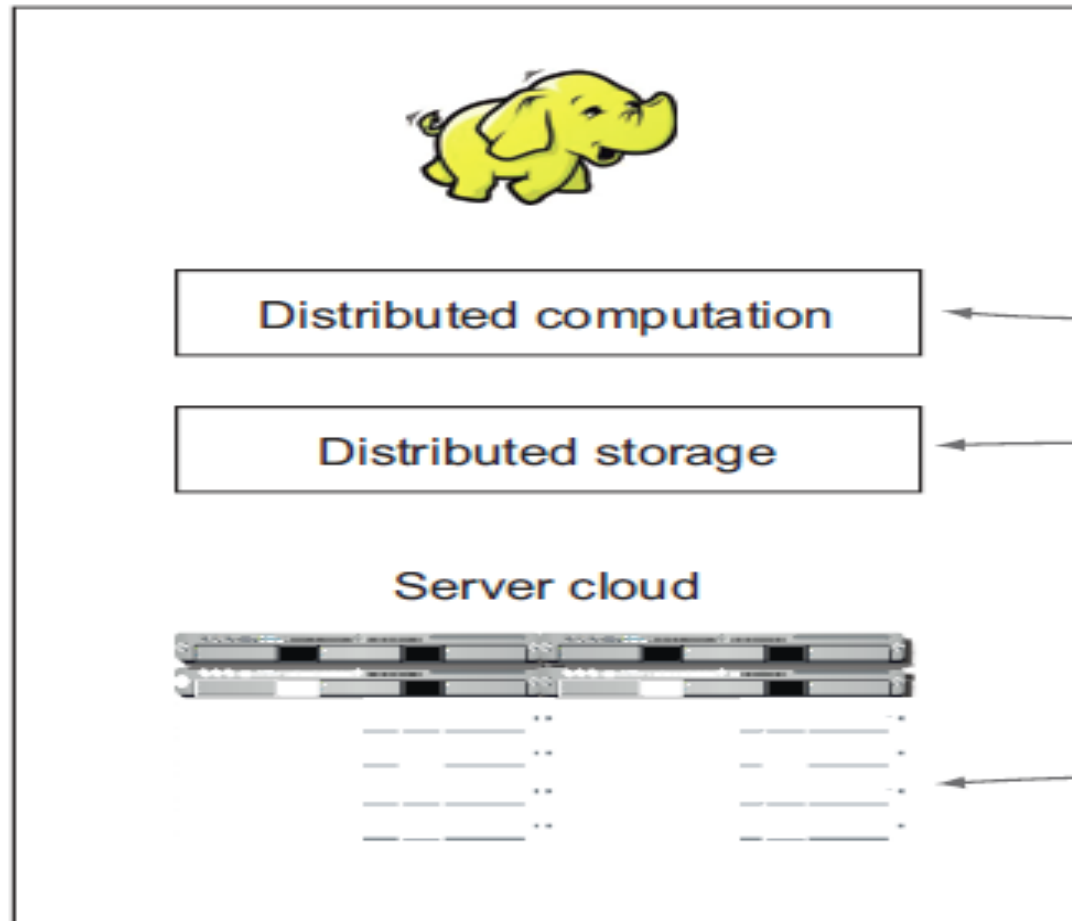
Designed to answer the question:

**“How to process big data with reasonable cost and time?”**

# layers



# hadoop environment - a distributed system

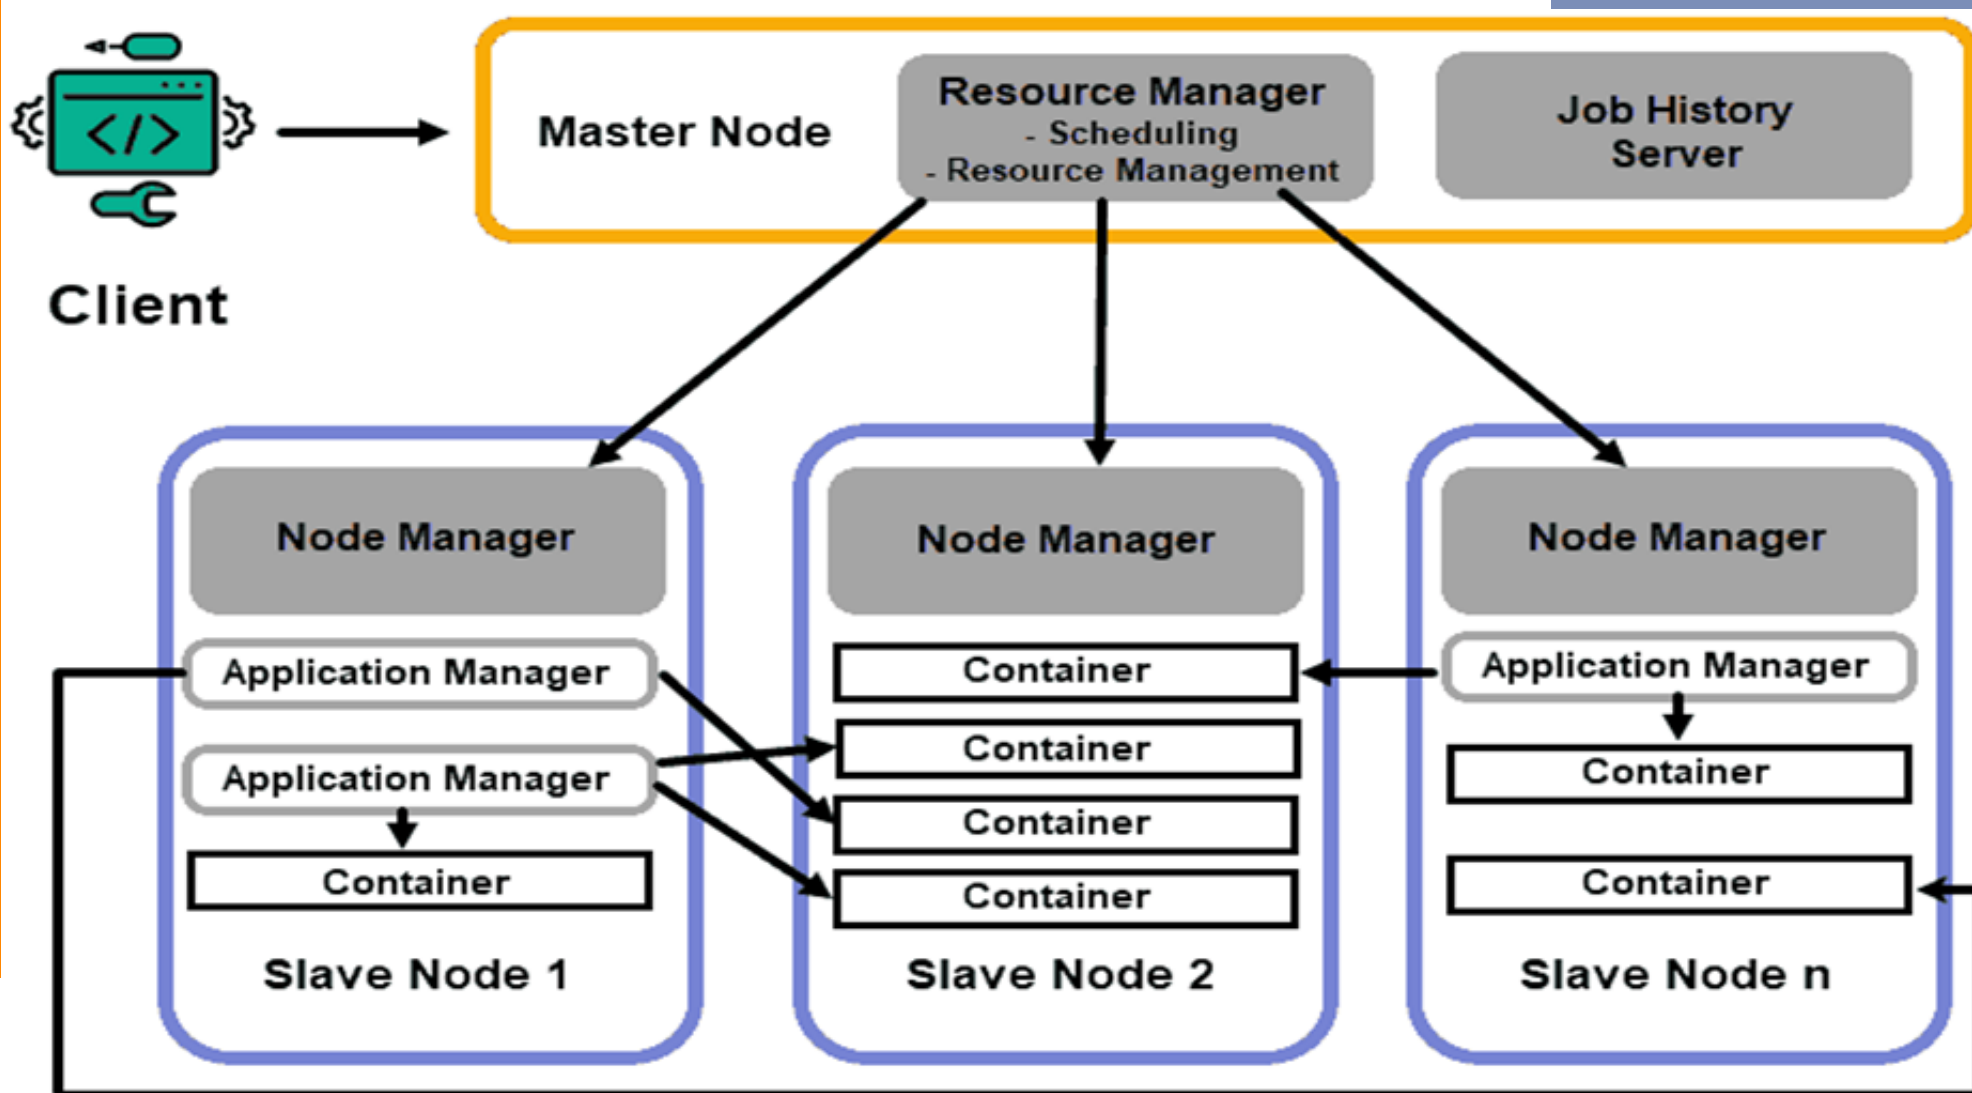


The computation tier is a general-purpose scheduler and a distributed processing framework called MapReduce.

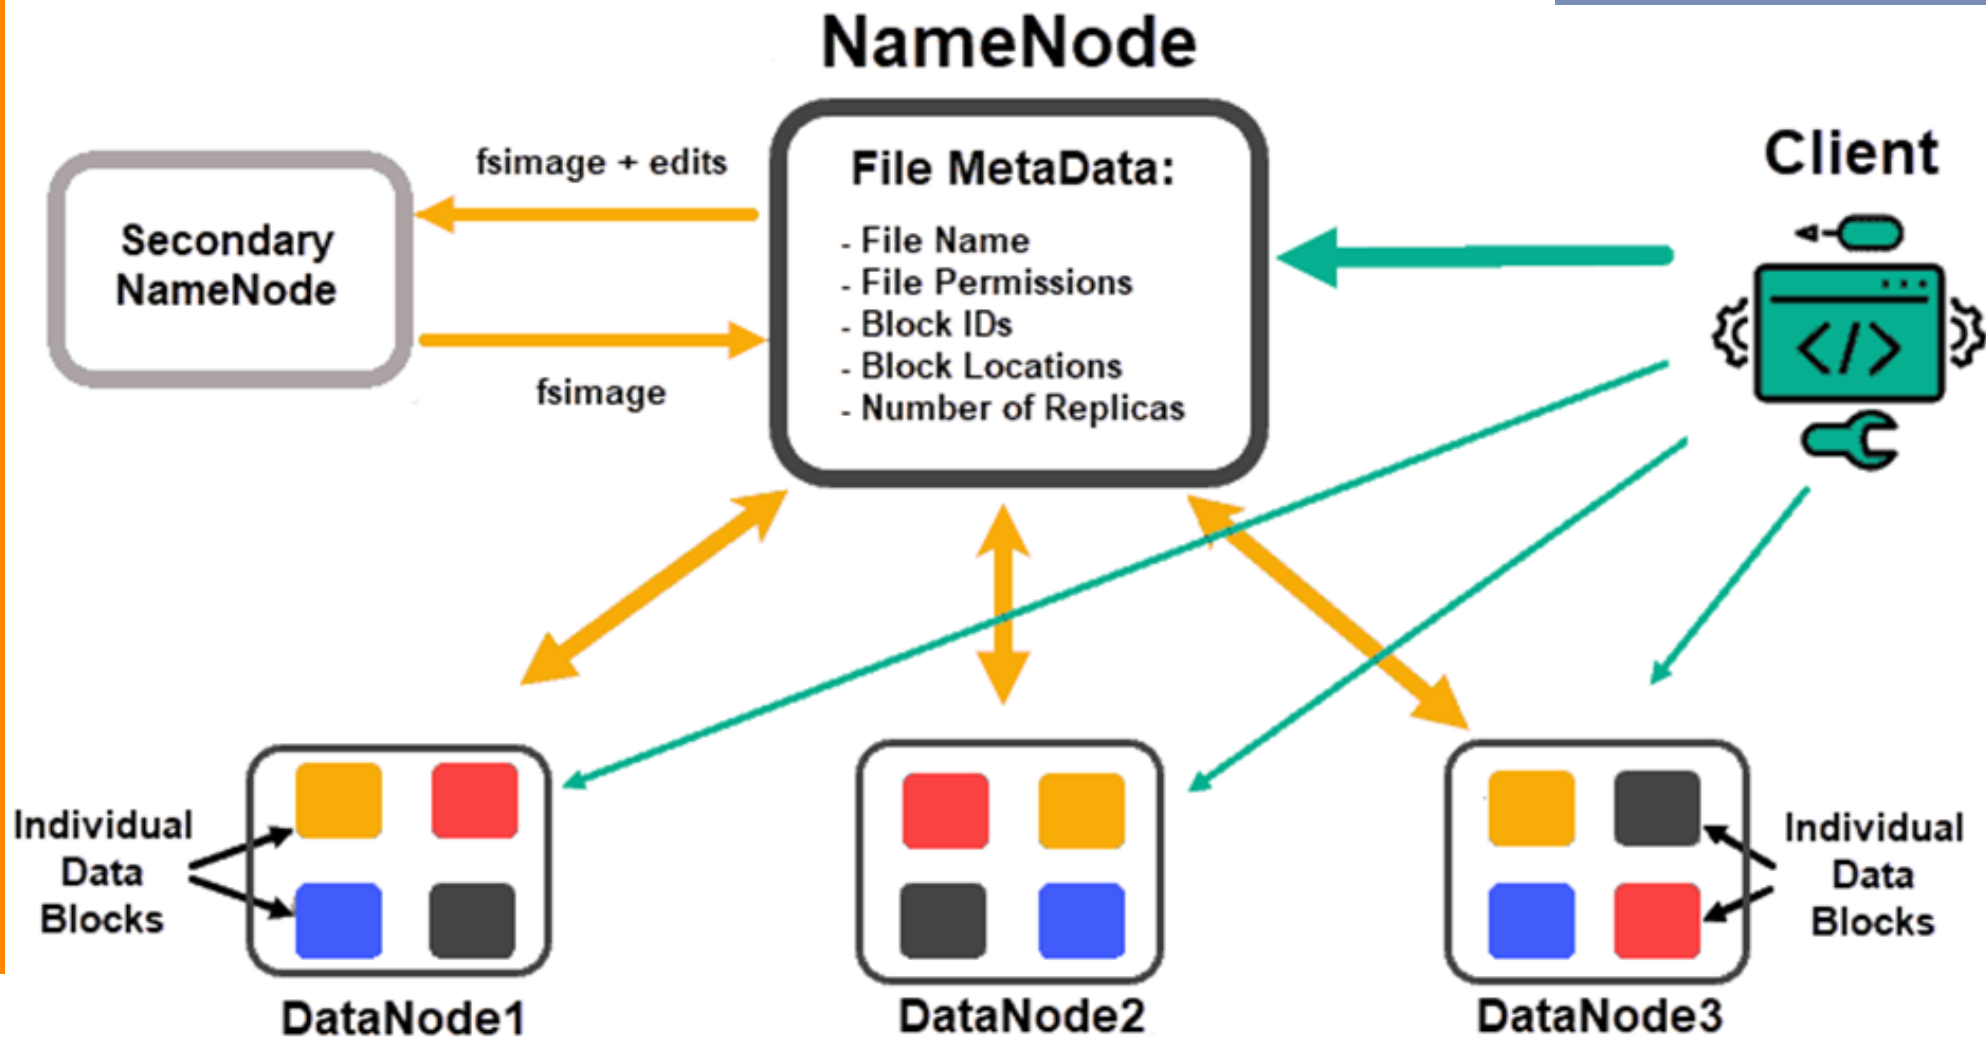
Storage is provided via a distributed filesystem called HDFS

Hadoop runs on commodity hardware.

# architecture

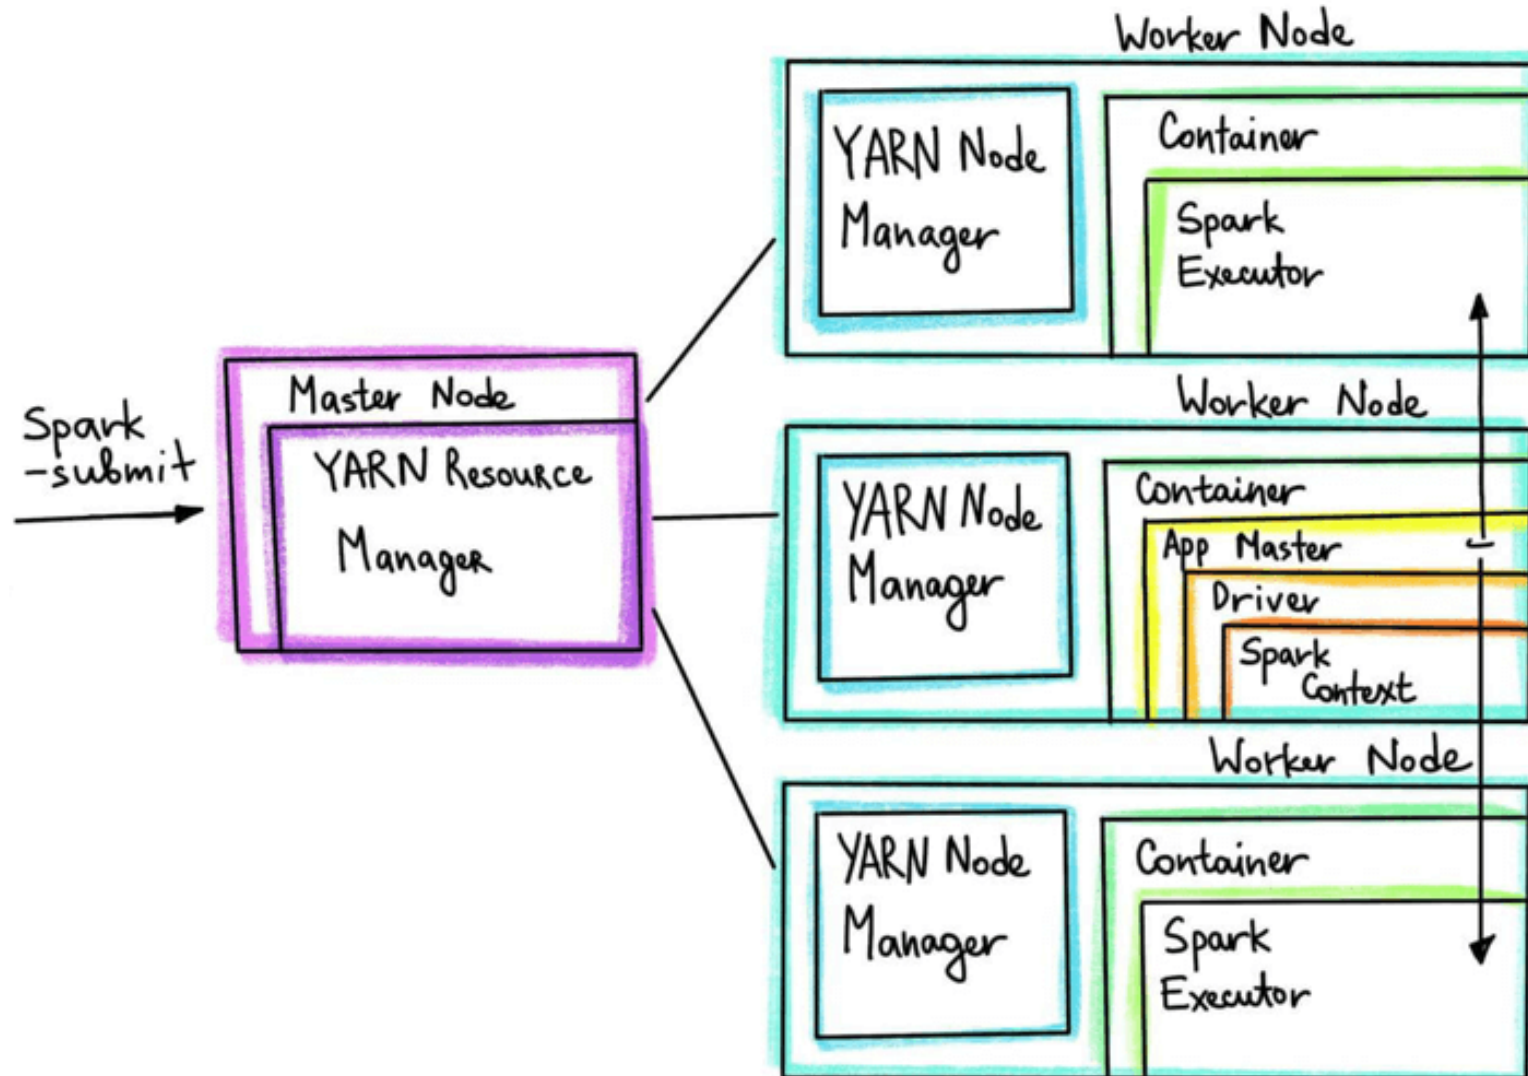


# hdfs



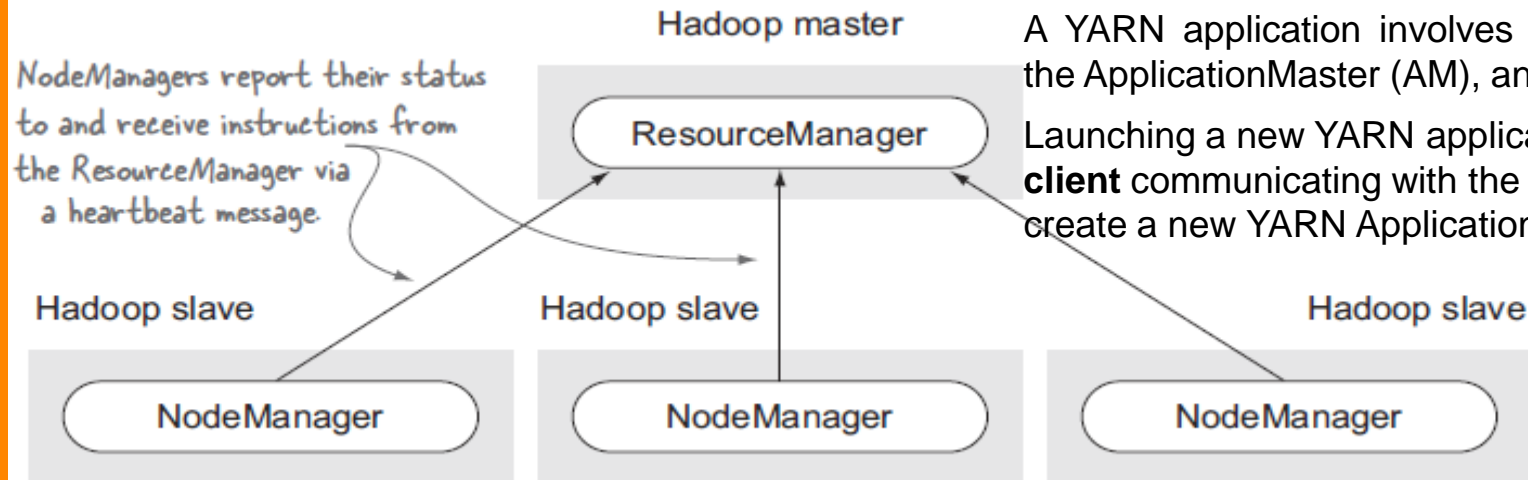


# yarn





# interactions of a yarn application

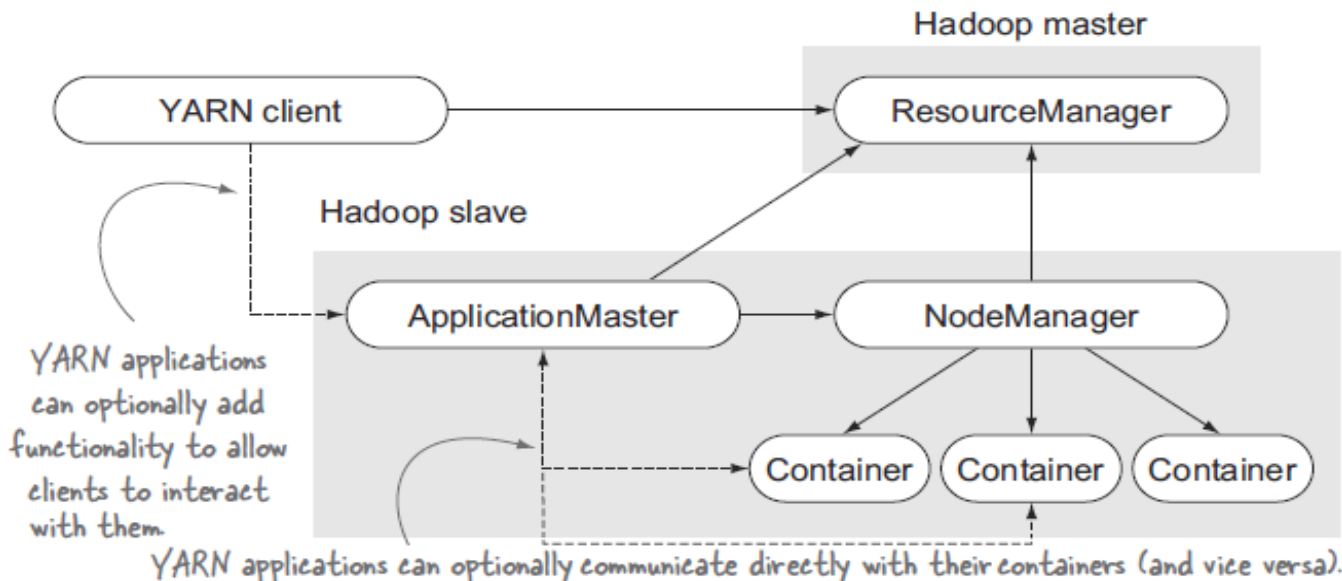


A YARN application involves 3 components—the client, the ApplicationMaster (AM), and the container.

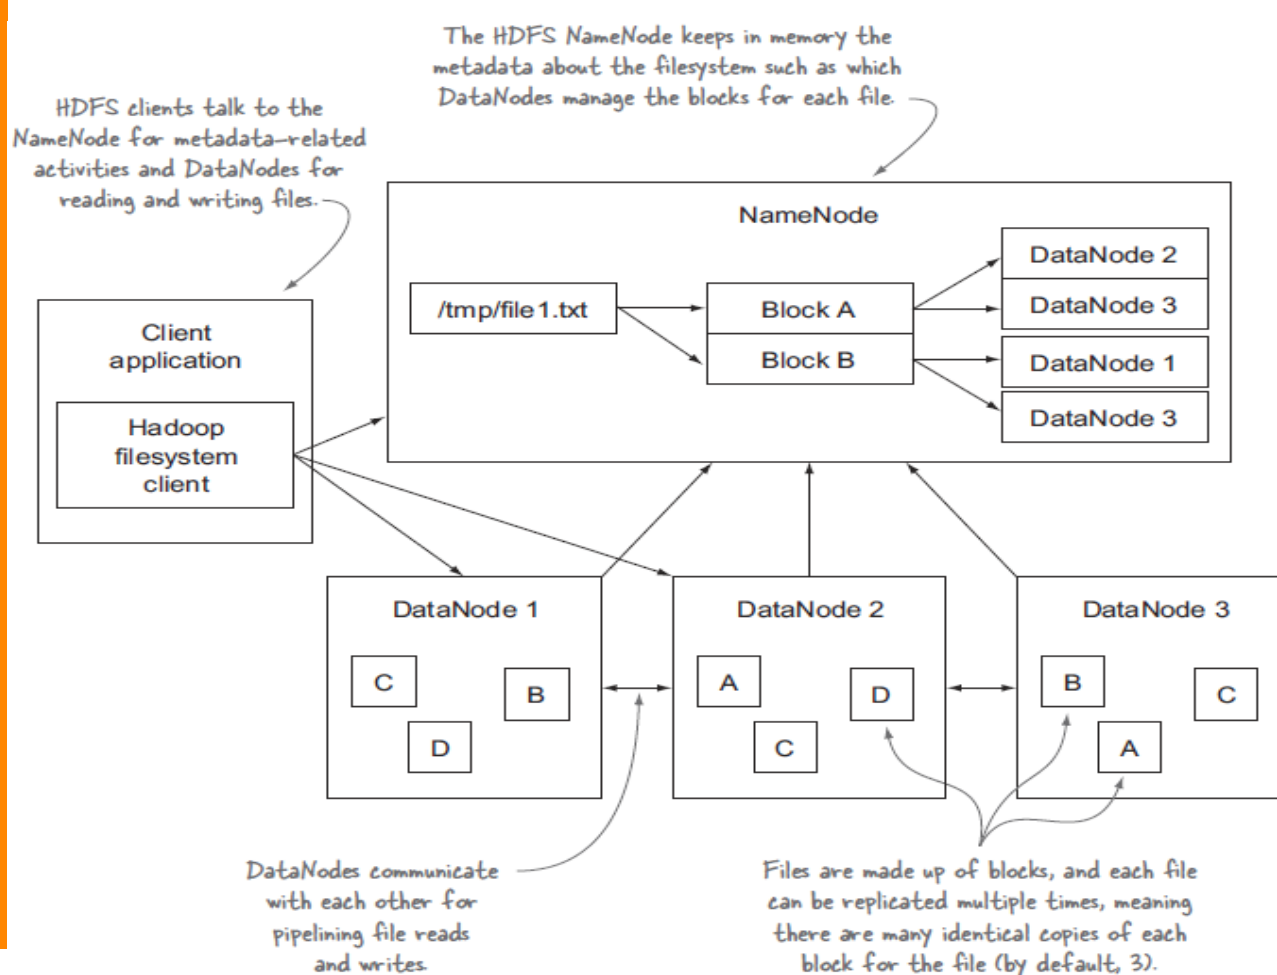
Launching a new YARN application starts with a YARN **client** communicating with the ResourceManager to create a new YARN ApplicationMaster instance.

the **ApplicationMaster** must specify the resources that each container requires in terms of which host should launch the container and what the container's memory and CPU requirements are.

A **container** is an application-specific process that's created by a NodeManager on behalf of an ApplicationMaster.




# hdfs client communication



**NameNode**  
File metadata:  
/user/chuck/data1 -> 1,2,3  
/user/james/data2 -> 4,5

**DataNodes**



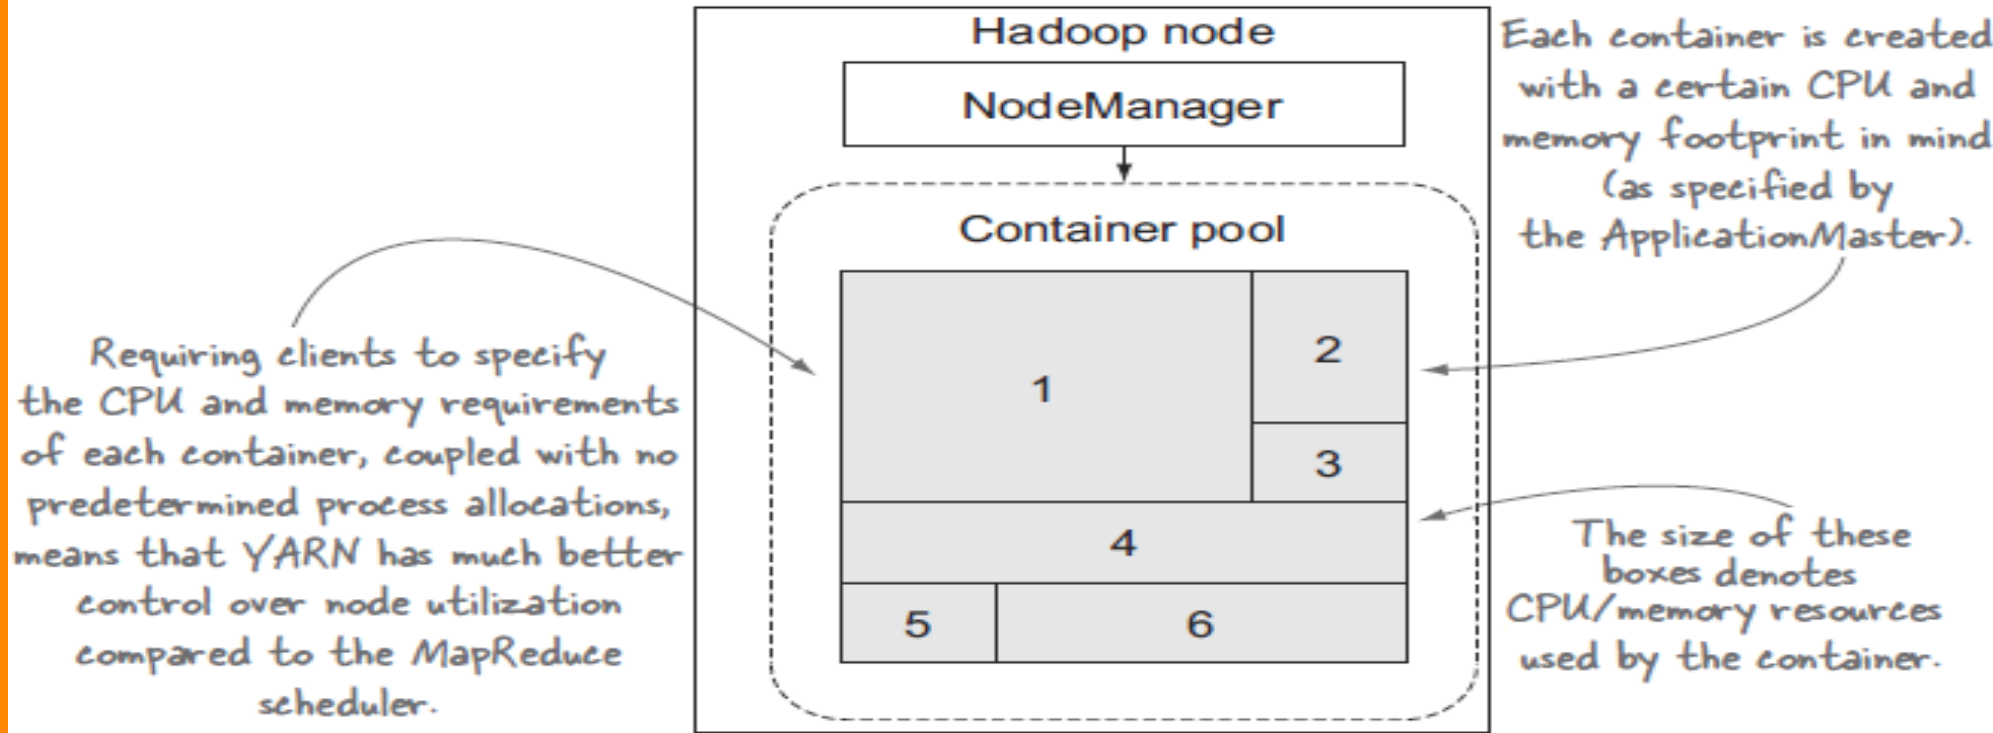
Hadoop 2 introduced two significant new features for HDFS—Federation and **High Availability (HA)**:

- Federation allows HDFS **metadata** to be **shared** across **multiple NameNode** hosts, which aides with HDFS scalability and also provides data isolation, allowing different applications or teams to run their own NameNodes without fear of impacting other NameNodes on the same cluster.

- High Availability in HDFS removes the single point of failure that existed in Hadoop 1, wherein a NameNode disaster would result in a cluster outage. HDFS HA also offers the ability for failover (the process by which a standby Name-Node takes over work from a failed primary NameNode) to be automated.

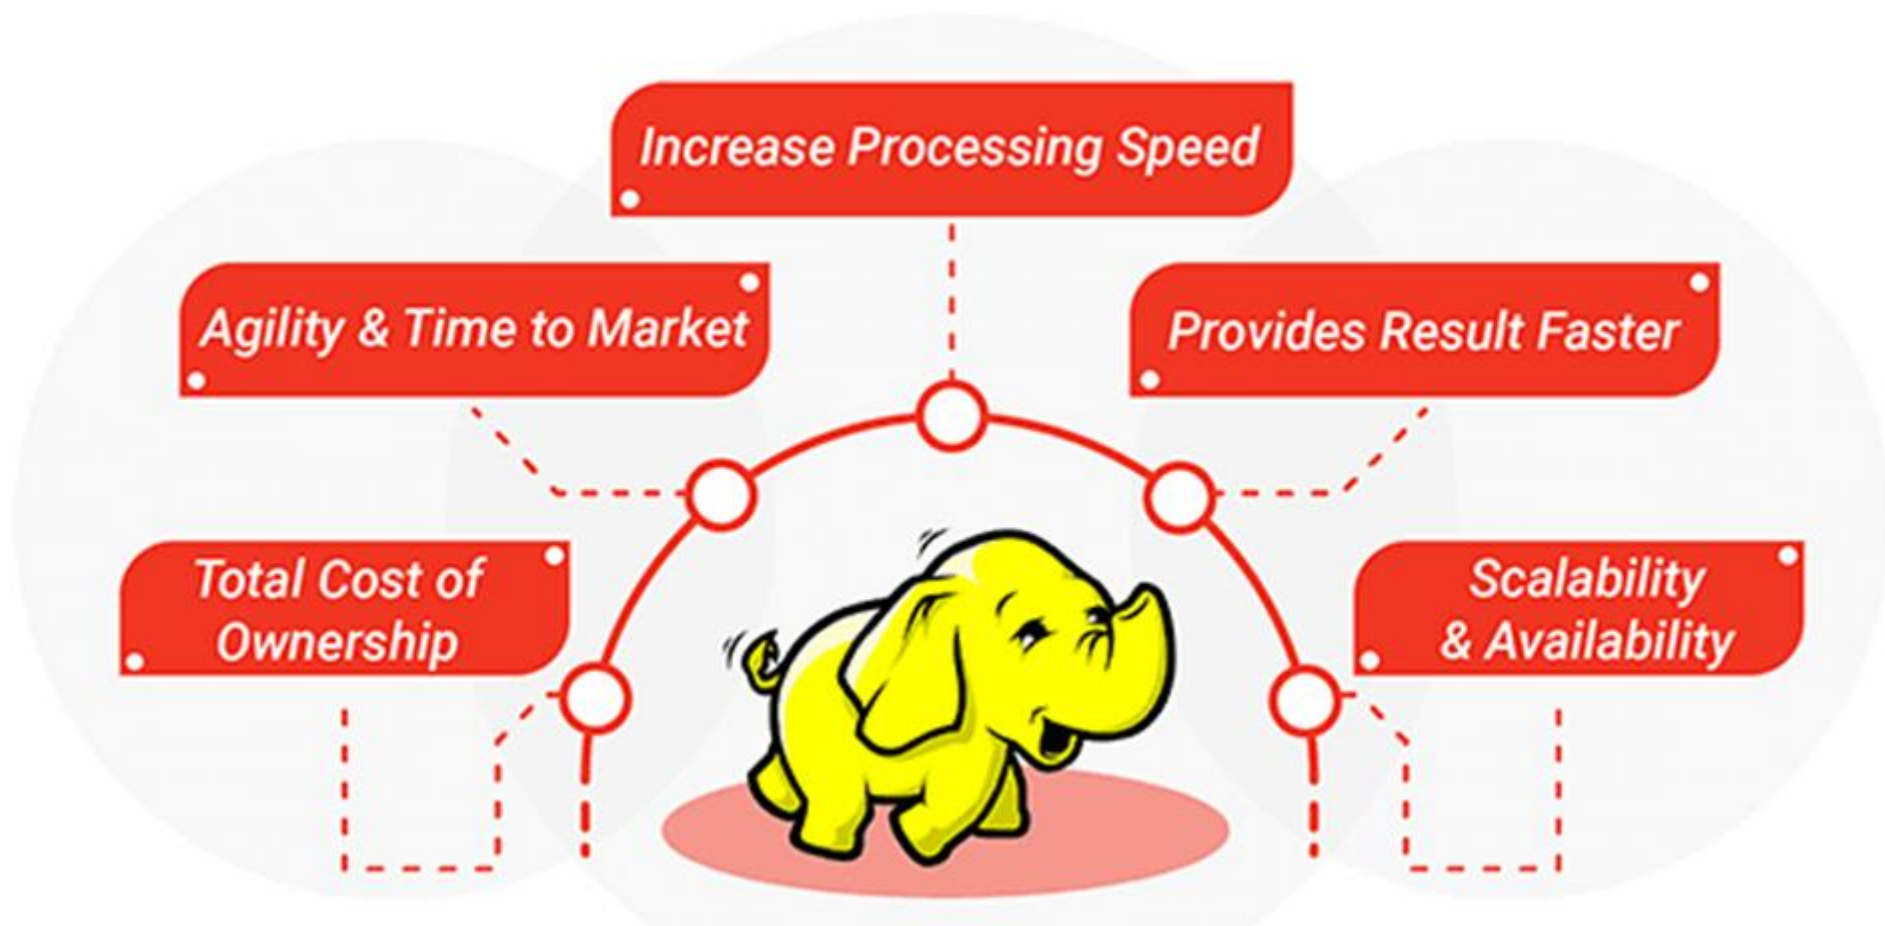
# Container configurations

The ability of the ResourceManager to schedule work based on exact resource requirements is a key to YARN's flexibility, and it enables hosts to run a mix of containers.




A container created by an ApplicationManager can be an arbitrary process—for example, a container process could simply be a Linux command such as `awk`, a Python application, or any process that can be launched by the operating system. This is the power of YARN—the ability to launch and manage any process across any node in a Hadoop cluster.


# hadoop with gpu




# benefits



## ***hadoop 3.0*** **BENEFITS**



Support GPUs.



Support  
Multiple  
Standby  
NameNodes



Supports  
multiple  
NameNodes  
for multiple  
namespaces.



Storage  
overhead  
reduced from  
200% to 50%.



Intra-Node  
Disk  
Balancing

