

ORACLE®

Certified Professional



Oracle Certified  
Professional Java  
Programmer

Microsoft Certified  
Professional

sohailimran@yahoo.com



سہیل عمران Sohail IMRAN 

*intro*


# intro

**MapReduce is a programming model or pattern within the Hadoop framework that is used to access big data stored in the Hadoop File System (HDFS).**

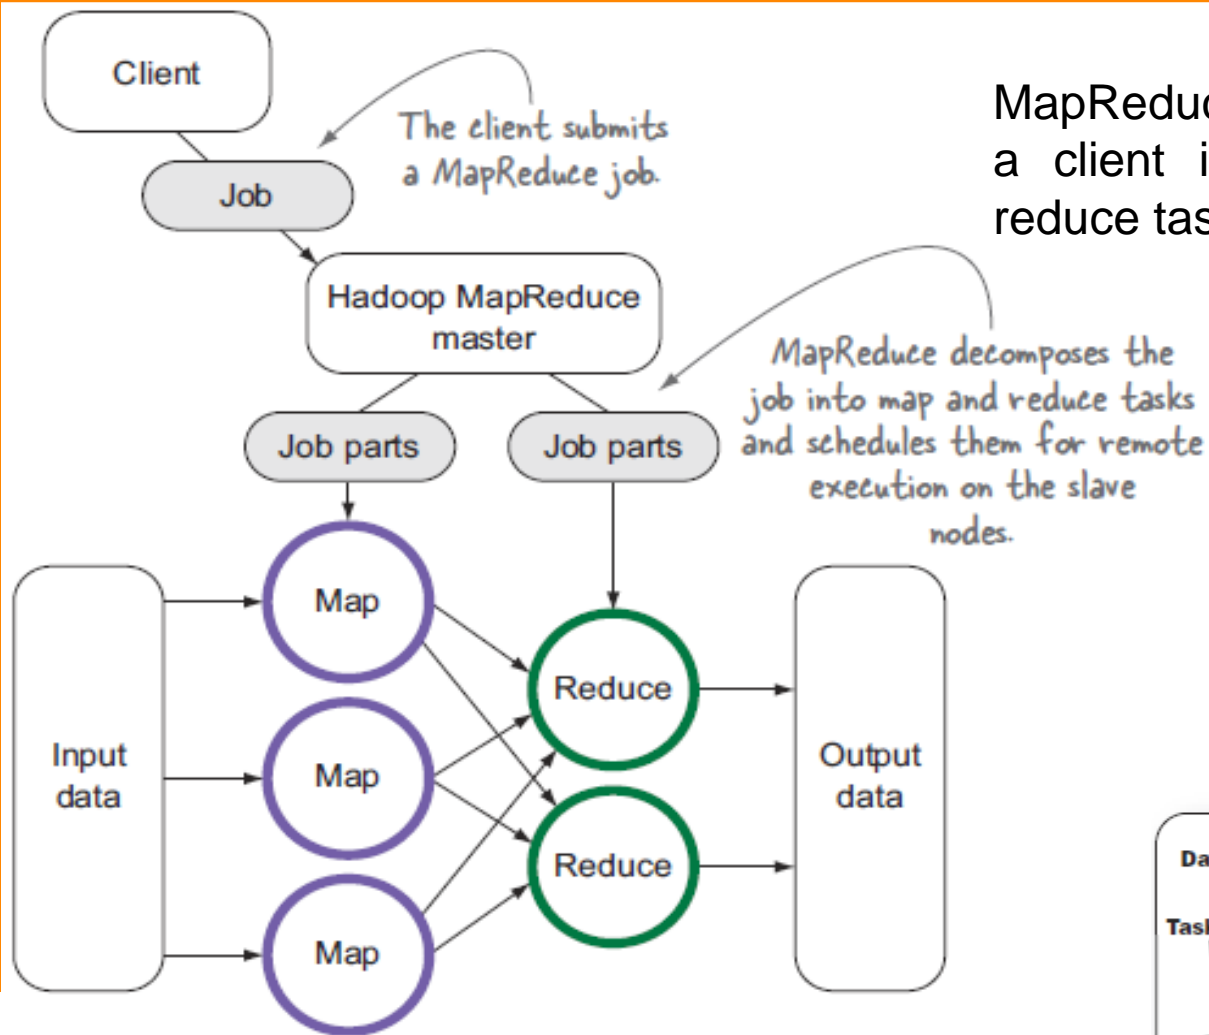
It is a core component, integral to the functioning of the Hadoop framework.

**With MapReduce, rather than sending data to where the application or logic resides, the logic is executed on the server where the data already resides, to expedite processing.**

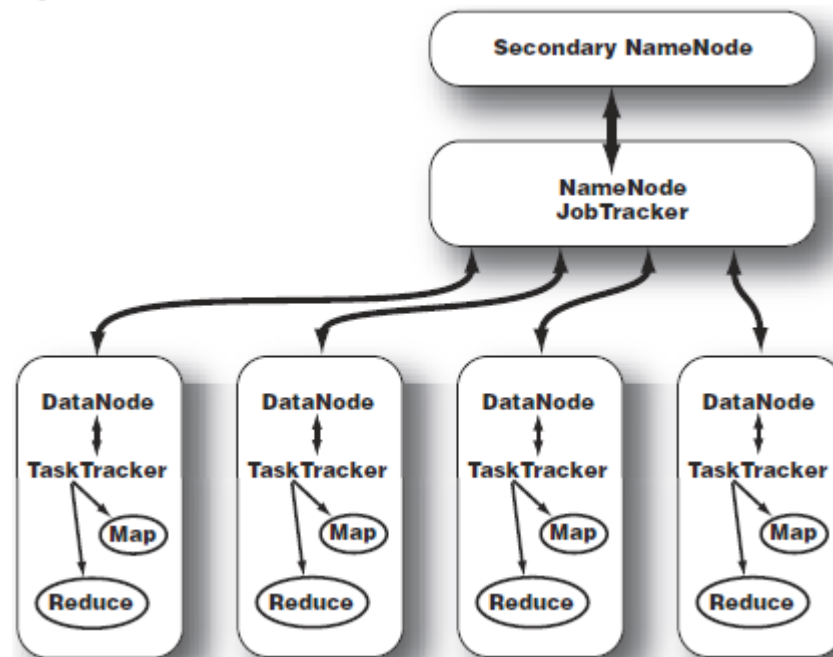
Data access and storage is disk-based—the input is usually stored as files containing structured, semi-structured, or unstructured data, and the output is also stored in files.



# the mapreduce model

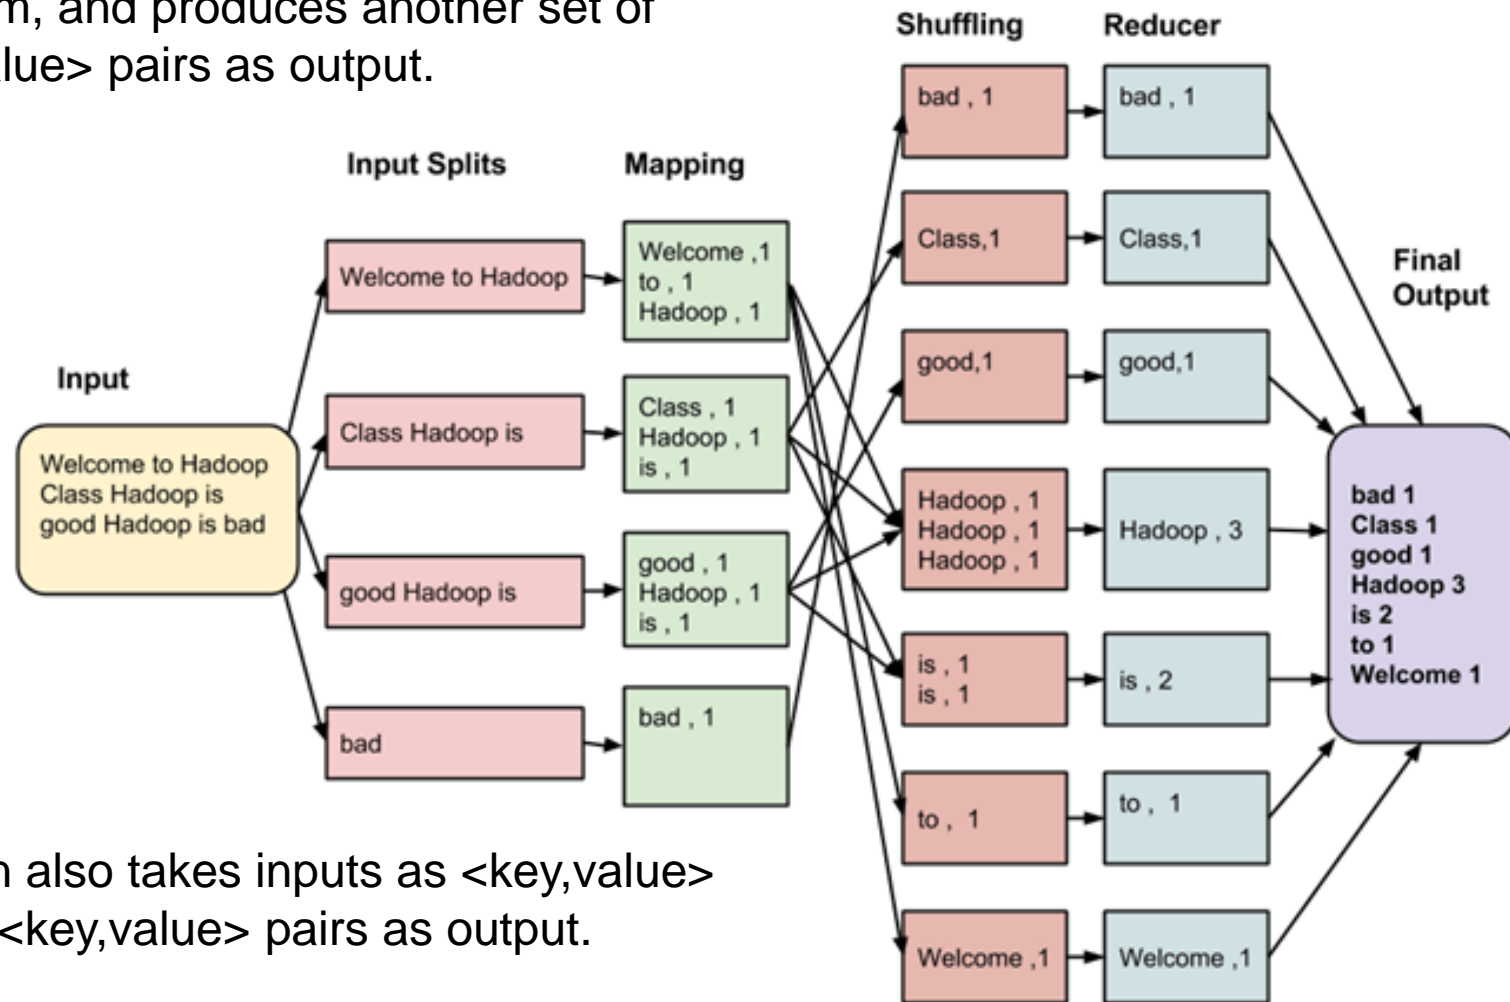


MapReduce decomposes work submitted by a client into small parallelized map and reduce tasks.



# example

The **Map** function takes input from the disk as <key,value> pairs, processes them, and produces another set of intermediate <key,value> pairs as output.



The **Reduce** function also takes inputs as <key,value> pairs, and produces <key,value> pairs as output.

# architecture

Hadoop divides the job into two types of tasks:

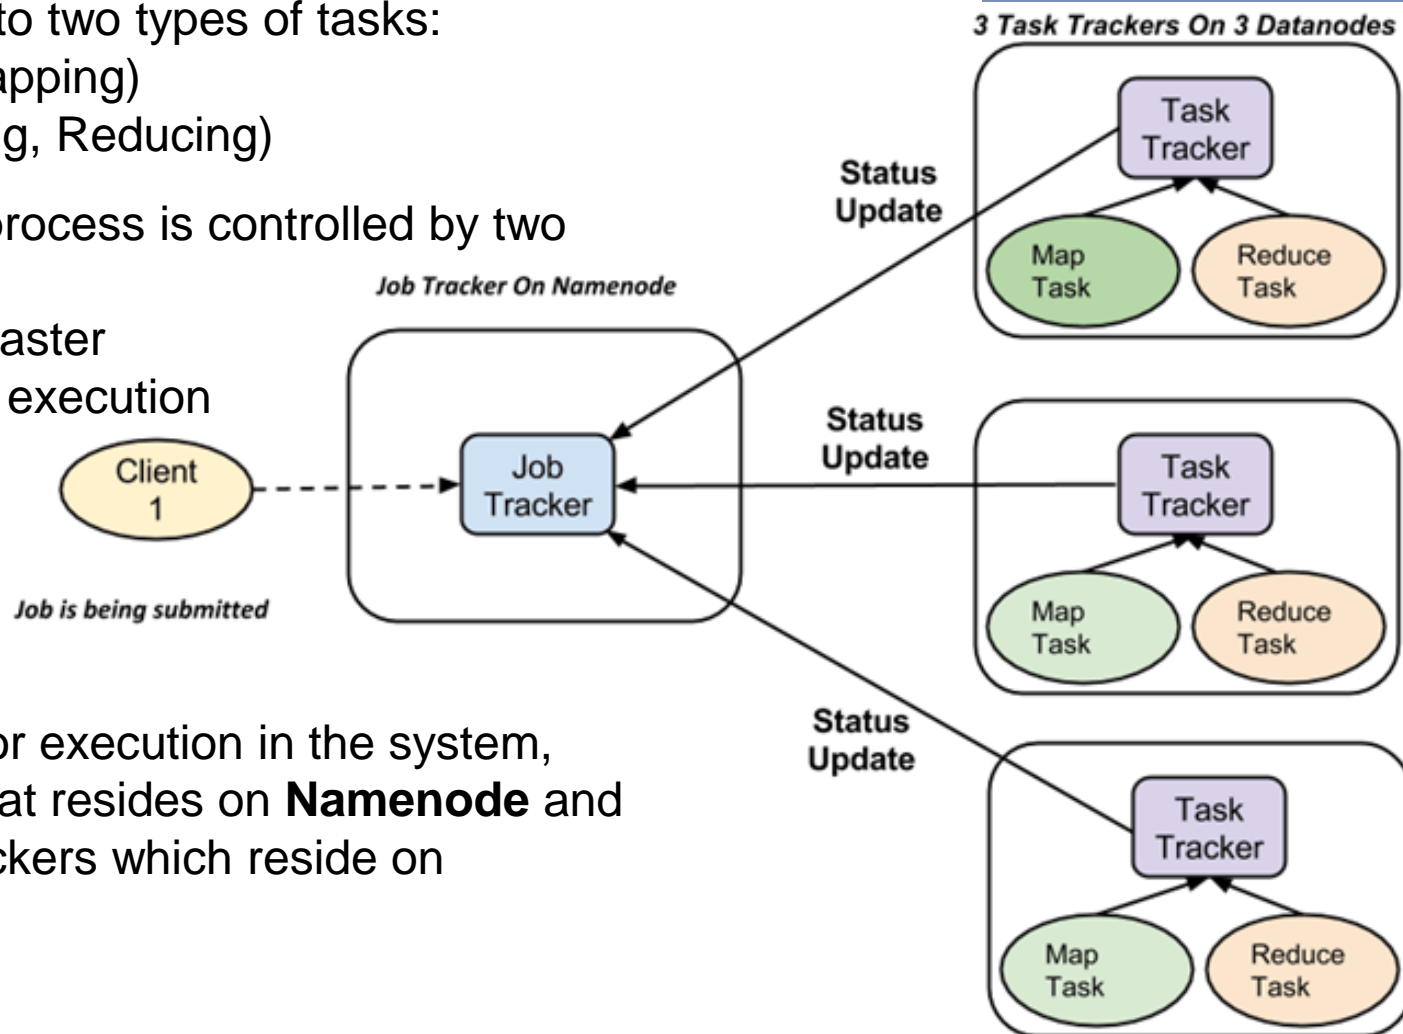
- **Map** tasks (Splits & Mapping)
- **Reduce** tasks (Shuffling, Reducing)

The complete execution process is controlled by two types of entities:

**Jobtracker:** Acts like a master (responsible for complete execution of submitted job)

**Multiple Task Trackers:**

Acts like slaves, each of them performing the job



For every job submitted for execution in the system, there is one Jobtracker that resides on **Namenode** and there are multiple tasktrackers which reside on **Datanode**.

# code

```
from pyspark.sql import SparkSession

# initialization of spark context
conf = SparkConf().setAppName(appName).setMaster(master)
sc = SparkSession\
    .builder\
    .appName("PythonWordCount")\
    .config(conf=conf)\
    .getOrCreate()

# read data from HDFS, as a result we get RDD of lines
linesRDD = sc.textFile("hdfs://...")

# from RDD of lines create RDD of lists of words
wordsRDD = linesRDD.flatMap(lambda line: line.split(" "))

# from RDD of lists of words make RDD of words tuples where
# the first element is a word and the second is counter, at the
# beginning it should be 1
wordCountRDD = wordsRDD.map(lambda word: (word, 1))

# combine elements with the same word value
resultRDD = wordCountRDD.reduceByKey(lambda a, b: a + b)

# write it back to HDFS
resultRDD.saveAsTextFile("hdfs://...")
spark.stop()
```